

✓ Honors Computational Genetics

Final Exam

- This exam is open notes but all of your answers should be your OWN work
- Partial credit will be given for ALL questions, so make sure to answer everything that you can
- Email Dr. Ellison with any questions: chris.ellison@rutgers.edu

When you finish

- Sign the honor pledge at the bottom of the notebook
- Make sure you completed the Canvas quiz
- Please **save a copy of the notebook to GitHub** BEFORE 8:00am on December 17th.

IMPORTANT

- *REMEMBER TO PERIODICALLY SAVE YOUR WORK!*
- *All PYTHON code questions worth 10 or more points should have comments*
- *DON'T FORGET TO DO THE MULTIPLE CHOICE QUESTIONS ON CANVAS UNDER THE QUIZZES TAB!*

```
# This code mounts your google drive to this notebook so that you can access the files in the google drive folder we shared with
from google.colab import drive
drive.mount('/content/drive')
!ln -s /content/drive/MyDrive/gen203/final_exam_data
```

Mounted at /content/drive

```
# This code installs a set of bioinformatics software packages that you will need for the exam
!pip install biopython
!pip install -q condacolab
import condacolab
condacolab.install_from_url("https://github.com/Ellison-Lab/gen203-condacolab-installer/releases/download/0.2.0/gen203_condacolab
```

```
Collecting biopython
  Downloading biopython-1.84-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (12 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from biopython) (1.26.4)
  Downloading biopython-1.84-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (3.2 MB)
3.2/3.2 MB 25.8 MB/s eta 0:00:00
Installing collected packages: biopython
Successfully installed biopython-1.84
Downloading https://github.com/Ellison-Lab/gen203-condacolab-installer/releases/download/0.2.0/gen203_condacolab-0.2.0-Li
Installing...
Adjusting configuration...
Patching environment...
Done in 0:01:12
Restarting kernel...
```

✓ 01

14 points

Use Unix commands to do the following:

- ✓ (A) make three new directories to produce the following path: /content/A/B/C

```
%bash
```

```
mkdir /content/A
mkdir /content/A/B
mkdir /content/A/B/C
```

bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)

- ✓ (B) Use a single command to move into the directory C and then use another Unix command that prints your current location

```
%%bash
```

```
cd /content/A/B/C  
pwd
```

```
↵ /content/A/B/C  
bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)
```

✓ (C) Use Unix commands to create a new file named `seq1.fasta`.

- The file should be in FASTA format with the sequence ID `seq1` and the DNA sequence `GATTACA`
- After making the file, use a Unix command to display the contents of the file on your screen

```
%%bash
```

```
echo '>seq1' > seq1.fasta  
echo 'GATTACA' >> seq1.fasta
```

```
cat /content/seq1.fasta
```

```
↵ >seq1  
GATTACA  
bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)
```

✓ (D) copy the file `seq1.fasta` to a new file named `seq1_copy.fasta`

```
!cp seq1.fasta seq1_copy.fasta
```

```
↵ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
```

✓ (E) Use a single Unix command to move `seq1.fasta` and `seq1_copy.fasta` so that they are located inside the directory `C`.

```
!mv seq1.fasta seq1_copy.fasta /content/A/B/C
```

```
↵ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
```

✓ (F) Use a single command to move into directory `C` and another command to list the contents of the directory.

```
%%bash
```

```
cd /content/A/B/C  
ls
```

```
↵ seq1_copy.fasta  
seq1.fasta  
bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)
```

✓ (G) Use a single command to move into directory `C` and then:

- Use a command to search `seq1.fasta` for the character that precedes the sequence ID in a FASTA file
- Use another command to print the contents of both FASTA files to the screen

```
%%bash
```

```
cd /content/A/B/C/  
grep '>' seq1.fasta  
cat seq1.fasta seq1_copy.fasta
```

```
↵ >seq1  
>seq1  
GATTACA  
>seq1  
GATTACA  
bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)
```

✓ 02

21 points

The file `exam_data/genome_features.tab` contains the coordinates for various types of genome features in the *Drosophila* genome. The tab-delimited columns in the file are:

1. Chromosome ID
2. Feature start coordinate
3. Feature end coordinate
4. Source of feature (e.g. the name of the group/program that annotated the feature)
5. Type of feature (e.g. what kind of feature it is)

Note: Although it may appear that the features in this file are sorted according to their start coordinates, this may not be true for all of the features and you will need to use Unix to ensure they are properly sorted for some of the questions.

Use Unix commands to answer the questions below:

- ✓ (A) How many lines are in the file?

```
!wc -l exam_data/genome_features.tab
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
85740 exam_data/genome_features.tab
```

- ✓ (B) Print the first 5 lines

```
!head -n 5 exam_data/genome_features.tab
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
2L      7874      8318      DRSC_dsRNA      RNAi_reagent
2L      8238      8534      BKN_dsRNA       RNAi_reagent
2L      8267      8580      DRSC_dsRNA      RNAi_reagent
2L      8267      8580      HFA_dsRNA       RNAi_reagent
2L      10964     10984     TriP_4          RNAi_reagent
```

- ✓ (C) Print the last 5 lines

```
!tail -n 5 exam_data/genome_features.tab
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
2R      25253314    25253440    FlyBase exon
2R      25253581    25253761    FlyBase exon
2R      25253820    25254078    FlyBase exon
2R      25254128    25254535    FlyBase exon
2R      25255009    25255318    FlyBase exon
```

- ✓ (D) Use a single command to print the number of features present on each chromosome that is listed in the file.

```
!cut -f 1,5 exam_data/genome_features.tab | sort | uniq -c
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
15623 2L      exon
24303 2L      RNAi_reagent
919 2L      transposable_element
17495 2R      exon
26042 2R      RNAi_reagent
1358 2R      transposable_element
```

- ✓ (E) Use a single command to print the first feature (i.e. the feature whose start coordinate is closest to zero) on chromosome 2L

```
!cut -f 1-5 exam_data/genome_features.tab | grep 2L | sort -k2,2 -n -r | tail -1 | column -t
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
2L 7529 8116 FlyBase exon
```

- ✓ (F) Use a single command to print the last feature (i.e. the feature with the largest end coordinate) on chromosome 2R

```
!cut -f 1-5 exam_data/genome_features.tab | grep 2R | sort -k2,3 -n -r | head -1 | column -t
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
```

```
2R 25272400 25272919 DRSC_dsRNA RNAi_reagent
```

- ✓ (G) Use a single command to print each type of feature whose source is from FlyBase, without duplicates

Note: feature type is listed in column number 5.

```
!cut -f 4,5 exam_data/genome_features.tab | grep FlyBase | sort | uniq
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
```

```
FlyBase exon
FlyBase transposable_element
```

- ✓ (H) How many features on chromosome 2L share the exact same start and end coordinates with at least one other feature?

- Use a single Unix command to answer this question

```
!cut -f 1,2,3 exam_data/genome_features.tab | grep 2L | sort | uniq -d | wc -l
```

```
➜ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
```

```
5438
```

✓ 03

15 points

Use the multiple sequence alignment in the file `exam_data/msa.fasta` to answer the questions below

- ✓ (A) use biopython to print the following information about the multiple sequence alignment:

- There are _____ sequences and _____ columns in the alignment.
- _____ columns contain only a single gap
- _____ columns contain more than one gap

```
from Bio import AlignIO

align = AlignIO.read('exam_data/msa.fasta', 'fasta')

# number of columns and sequences
alen = len(align[0])
num_seq = 0
for i in align:
    num_seq += 1
# print(num_seq)

# counters for single and multiple gaps
single_gap_col = 0
multiple_gap_col = 0

# iterate through each column
for j in range(alen):
    column = align[:,j]
    gaps = column.count('-') # count number of gaps
    # print(gaps)

    if gaps == 1:
        single_gap_col += 1
    elif gaps > 1:
        multiple_gap_col += 1
    else:
        continue

print("There are", num_seq, "sequences and", alen, "columns in the alignment.")
```

```
print(single_gap_col, "columns contain only a single gap")
print(multiple_gap_col, "columns contain more than one gap")
```

```
→ There are 4 sequences and 610 columns in the alignment.
    11 columns contain only a single gap
    15 columns contain more than one gap
```

✓ (B) Use biopython to make a distance matrix for this alignment

```
from Bio.Phylo.TreeConstruction import DistanceCalculator
from Bio import AlignIO
```

```
aln = AlignIO.read('exam_data/msa.fasta', 'fasta')
calculator = DistanceCalculator('identity')
dm = calculator.get_distance(aln)
print(dm)
```

```
→ sequence1 0.000000
sequence2 0.024590 0.000000
sequence3 0.013115 0.037705 0.000000
sequence4 0.047541 0.022951 0.060656 0.000000
    sequence1 sequence2 sequence3 sequence4
```

✓ (C) What two sequences are the most similar?

sequence1 and sequence3

✓ (D) What two sequences are the most dissimilar?

sequence3 and sequence4

✓ 04

10 points

✓ Use biopython to parse the Illumina sequences in exam_data/yeast.fastq.

- Write the sequences whose average quality score is 35 or greater in FASTQ format to the file my_filtered.fastq.
- Print the total number of sequences in the original file, the number of sequences in the filtered file, and the number of sequences that were discarded:

Total Sequences: _____

High Quality Sequences: _____

Discarded Sequences: _____

```
from Bio import SeqIO
import statistics
```

```
total = 0
high_quality = 0
discarded = 0
```

```
filtered = open('my_filtered.fastq', 'w')
```

```
# iterate over the original FASTQ file
for record in SeqIO.parse('exam_data/yeast.fastq', 'fastq'):
    total += 1
```

```
    # find quality scores for each sequence
    scores = record.letter_annotations["phred_quality"]
    avg_score = statistics.mean(scores)
```

```
    # filtering based on avg score of 35
```

```
    if avg_score >= 35:
```

```
        high_quality += 1
```

```
        SeqIO.write(record, filtered, 'fastq') # writing record to new file
```

```
        seqio.write(record, filtered, fastq)    # writing scores to new file
    else:
        discarded += 1

print("Total Sequences:", total)
print("High Quality Sequences:", high_quality)
print("Discarded Sequences:", discarded)

filtered.close()    # I didn't forget this time :)
```

```
↵ Total Sequences: 1410
   High Quality Sequences: 1393
   Discarded Sequences: 17
```

▼ 05

7 points

▼ Use the file `exam_data/yeast.fastq` for this question

- Align these sequences to the yeast chromosome #8, which is in the file `exam_data/chr08.fasta`
- Convert your alignments to a sorted BAM file named `my_sorted_alignments.bam`

%%bash

```
bowtie2-build exam_data/chr08.fasta chr08.fasta
bowtie2 --no-unal -x chr08.fasta -U exam_data/yeast.fastq -S my_sorted_alignments.sam
samtools view -b my_sorted_alignments.sam > my_sorted_alignments.bam
```

```
↵
```

```

Renaming chr08.fasta.1.bt2.tmp to chr08.fasta.1.bt2
Renaming chr08.fasta.2.bt2.tmp to chr08.fasta.2.bt2
Renaming chr08.fasta.rev.1.bt2.tmp to chr08.fasta.rev.1.bt2
Renaming chr08.fasta.rev.2.bt2.tmp to chr08.fasta.rev.2.bt2
1410 reads; of these:
  1410 (100.00%) were unpaired; of these:
    0 (0.00%) aligned 0 times
    1410 (100.00%) aligned exactly 1 time
    0 (0.00%) aligned >1 times
100.00% overall alignment rate

```

Write a short paragraph addressing the points below:

- What information is contained in the SAM file?
- How is the SAM format different from the BAM file format?
- Why would you want to convert a SAM file to a BAM file?
- What gets sorted when you sort a BAM file and in what order?

The Sequence Alignment/Map (SAM) stores the Illumina sequence ID and quality scores of the alignment, the alignment location including the reference sequence ID and position mapping quality, the locations of mismatches and/or gaps in alignment, and the alignment strand, whether it is in forward or reverse. The SAM format is tab-delimited; the alignments have to be compressed to BAM format, which is only in binary and, therefore, not readable by humans. By compressing into a BAM format, the software can sort the alignments based on the genome positions, aided by the compressed format, which saves space.

✓ 06

9 points

The yeast chromosome 8 is over 500 kb long but the Illumina sequences in `exam_data/yeast.fastq` are from a much smaller contiguous region, located somewhere on the chromosome.

IMPORTANT: these questions purposefully do not tell you whether to use Unix or Python, nor do they tell you whether there is a specific Unix program that will accomplish the task. However, you have done all of these things previously in the homework assignments.

(A) Create a tab-delimited file that lists the per-base sequencing coverage for each position of chromosome 8

```

%%bash
samtools sort my_sorted_alignments.bam > my_really_sorted_alignments.bam
samtools depth -a my_really_sorted_alignments.bam > coverage_per_position.tab

```

bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)

✓ (B) Determine the start and end positions of the chromosome 8 region that was sequenced.

Print the start and end positions, assuming the first nucleotide of chromosome 8 is position 1.

```

read_coverage = open("coverage_per_position.tab")
position = []

for coverage in read_coverage:
    coverage = coverage.rstrip()
    cov_list = coverage.split("\t")

    num_pos = int(cov_list[1])
    depth = int(cov_list[2])

    if depth > 0:
        position.append(num_pos)

min_value = min(position)
max_value = max(position)

print("Read Coverage begins at position", min_value)
print("Read Coverage ends at position", max_value)

read_coverage.close()

```

Read Coverage begins at position 28166
Read Coverage ends at position 32936

✓ (C) All of the Illumina sequences in `exam_data/yeast.fastq` are 100 bp long.

Calculate the average sequencing coverage for the region from (B) above.

Show your work.

```
num_reads = 1410    # from bowtie2
size = 32936 - 28166
length = 100
```

```
coverage = (num_reads * length)/size
print(coverage)
```

```
→ 29.559748427672957
```

✓ 07

9 points

(A) Use Velvet to assemble the Illumina sequences in `exam_data/yeast.fastq` by using a Kmer size of 27

```
!velveth YEAST 27 -short -fastq exam_data/yeast.fastq
```

```
→ /bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
[0.000000] Reading FastQ file exam_data/yeast.fastq;
[0.004120] 1410 sequences found
[0.004139] Done
[0.066690] Reading read set file YEAST/Sequences;
[0.067218] 1410 sequences found
[0.075104] Done
[0.075119] 1410 sequences in total.
[0.075286] Writing into roadmap file YEAST/Roadmaps...
[0.077499] Inputting sequences...
[0.077621] Inputting sequence 0 / 1410
[0.197294] === Sequences loaded in 0.119801 s
[0.197349] Done inputting sequences
[0.197357] Destroying splay table
[0.200392] Splay table destroyed
```

```
!velvetg YEAST -exp_cov 30 -cov_cutoff auto
```

```
→
```



```

[0.243937] == Nodes Scaffolded in 0.033618 s
[0.243985] Preparing to correct graph with cutoff 0.200000
[0.244062] Cleaning memory
[0.244088] Deactivating local correction settings
[0.244106] Pebble done.
[0.244112] Concatenation...
[0.244118] Renumbering nodes
[0.244123] Initial node count 1
[0.244129] Removed 0 null nodes
[0.244134] Concatenation over!
[0.244139] Removing reference contigs with coverage < 10.957038...
[0.244149] Concatenation...
[0.244155] Renumbering nodes
[0.244160] Initial node count 1
[0.244165] Removed 0 null nodes
[0.244170] Concatenation over!
[0.244385] Writing contigs into YEAST/contigs.fa...
[0.245472] Writing into stats file YEAST/stats.txt...
[0.245656] Writing into graph file YEAST/LastGraph...
[0.247669] Estimated Coverage cutoff = 10.957038
Final graph has 1 nodes and n50 of 4745, max 4745, total 4745, using 1410/1410 reads

```

▼ (B) Use Unix to print your assembled sequence to the screen

```
!cat YEAST/contigs.fa
```



```

AATATAAAACAAAAACGATAACGGCGACTGTCAAAAAGCGAGGAGTTATTACGGTAATGA
GAAATATAATTGATGCGCATTGAATTAGGCAAAATATAGTTACTTCTAAGTATGGAATCA
ATTCTTGATCAACACCTTCGATGTCCTTTGAAAAGCGATTGATGTTCTACCAACCGGCG
TCACGTCGAAAAAATCGTATTTGGGCGATGTAGAACTAGATCTAGCAGATTATTAAGATCT
TTCTGGAGGCTCGCATACCGGATAAAAACGTCATCATTGTTTTAAACACCTAGCATTG
CCTGAATGATACCAATAAGAAAATATACGGTTAAGTAATAAAATGCATTATGTTTTATT
TCGAAGAGTCGGTCATCCCTTTAATGGCAGCGTGTCATCGCAAAACCTGGAGCATTTA
TTCGTACATTGGTATCGTTGACCCAATGTCGTATCCACCAAGACTGACTGATGAACAAAA
TTTGAGCTGTGATATAAAGAGCGAACAGGGCTGTTAAAGCTTTGAAGCTTCAAAAAATT
TCAGGTACCATTTATAAACATCGGGCTTATGGCACCGTTTGATTTTTCTTCTCTTCTA
TTAGCTGGCCATCATTGACAAAATTTGATCAAAAATTTATGTTCTCGGTGACAGGTTTGA
TTTTCTGAGAGTCATTTTTCTGGGAGCTTTAATCTATTAGCGTTCTTTTCATTAATGC
TATCTCGAGAAGAAAGTTGAACATATTTTTCTTAAAAAGCCCTTTGCTTTGTAATTCG
TAATAGTTCTTGATTCTTCACTTTGCCATTTTCCAACACAATCGCGAAATGGGCATTTT
TAAGTGTTAATGAAACATTGTGCGTAACATAAATGCAGGTTCTATTTTTTCATTAGTGGAC
CTGTGATGCAATTTTCATAGATCCATACAGCAGTATGTGAATCGACAGCGCTCAAAACAT
CATCTAGTAAGACATGCTTAGCACTCGAATAAACAGCTCTCGCCAAGGAAATTCCTGTT
TCTGCCCTCCTGATAAAGTTATACCCTTTTACCAATATCTGTTAGGTCACCTGCTGGTA
AAATCTCCAGGTCTCTTTTCAAGCCACATGCATCAATTACTTTGTTGTACCTATCCTCGT
TATAGAAGTTATCAAGATAATATTGTTTTTACCGTGTCATTTAATAGCCACGCACTTT
GTGAACAATATGCGAAGGAATTGGTTAAACCTTCGCAGTCGGGAATTAATCATGCTTTG
GTTCTAAGCTCGGAACAATGATAGAGCCACTAATTAGATTTAGTTCACCCAGTAAACCA
GCAGCAATGCACTTTTACCAGATCCTGTGAACCCAAAATCAAAATTTAACTTACCAATTT
GAAATTTAATATTCAAACACATAATTTGAATGCATTATATCGCTGTCAATTTTCATTCC
AGGTAAAGTCGCATTTTAAATCAATTTTATTTTGTCTGGAGATATGGTTAGTTGAT
TATATTTTTCTGTATCGTCCATCCTTAAAAAATCGCTTATCTTTTAGAGAGACTTTTG
ATTGATTTATGAACTTAGCATATTTGATAATTGATCCAGGGGTGCTTTAACAAAGTGA
AGAGTGACAAAGTAGTGAAAGCAAGCGGGCATTCAAATCTTCATGTTGAACAAATGTAC
AGATGGCGAAAGTGACACCTGTACCAAGGTCGGTGTCACGAACCAAGAAAAGAAGTTA
CGGACCACACCAAGATTTTTTAATAAGGATCTTAATTCCTTTTGCTTTATTGATTTGA
TTTCATTTATAATATTTCTTTCCCAAGCAAAATATTTGACAATCTTATGTTCTGTAAGC
ACTCGTTCAATTTTGAGATTTCTTTGGTCAGTACATTTCAAGTGTGCTTTTGAACTTAC

```

```
AAATAGTAGCTCAATAGCTGAGTTCCTCTAAACGAGTAAATCATCAACTTCCAAAAATG
AACCATTATTTTCGGATAATAAGGGGATCCgt
```

- ✓ (C) Is your assembled sequence a contig or a scaffold? Explain how you know

This is a single contig because there is one reported contig and the lack of N's indicative of scaffolds.

- ✓ (D) What is the best match for the sequence above in the Genbank refseq_protein database? Write a sentence listing the best match and explaining what you did to figure this out.

putative ATP-binding cassette multidrug transporter VMR1 [Saccharomyces cerevisiae S288C]

By using the NCBI BLASTX tool for a nucleotide sequence to protein, the results of 99.94% percent identity with an e-value of nearly 0 and a query coverage of 99% suggests this is the best match.

- ✓ 08

15 points

This question involves a different yeast gene that is 186 bp long

(A) Use BLAST to search the DNA sequence of the gene in exam_data/yeast_gene.fasta against the yeast chromosome 8 (exam_data/chr08.fasta).

- Use an e-value threshold of 1e-10
- Output the results in a tab-delimited format and let them print to the screen

```
%%bash
```

```
cp exam_data/chr08.fasta /content
makeblastdb -dbtype nucl -in chr08.fasta
blastn -query exam_data/yeast_gene.fasta -db chr08.fasta -outfmt 6 -evalue 1e-10 > output.tab
# cat output.tab
```



```
Building a new DB, current time: 12/17/2024 03:19:46
New DB name: /content/chr08.fasta
New DB title: chr08.fasta
Sequence type: Nucleotide
Keep MBits: T
Maximum file size: 10000000000B
Adding sequences from FASTA; added 1 sequences in 0.0115981 seconds.
bash: /usr/local/lib/libtinfo.so.6: no version information available (required by bash)
```

```
!cat output.tab
```



```
/bin/bash: /usr/local/lib/libtinfo.so.6: no version information available (required by /bin/bash)
YHR055C chrVIII 100.000 186 0 0 1 186 215064 214879 1.74e-96 344
YHR055C chrVIII 97.849 186 4 0 1 186 5601 5786 8.17e-90 322
```

- ✓ (B) How many copies of this gene are on chromosome 8?

Explain how you were able to tell how many copies there are and whether the copy/copies are likely to be real copies of the gene or just spurious results.

There are likely 2 real copies of the gene on chromosome 8 based on the high percent identity (100.00% and 97.849%) in consideration of a very low e-value.

- ✓ (C) Looking at your BLAST results, fill in the table below that lists the start and end coordinates on chrVIII where each copy of the gene is located.

Note that the start coordinate should refer to where the gene begins (i.e. its start codon) and the end coordinate should refer to where the gene ends (i.e. its stop codon).

Be sure to fill in the column in your table that lists whether the gene copy is on the forward or reverse strand of chromosome 8

Copy	Start	End	Strand
YHR055C	215064	214879	reverse
YHR055C	5601	5786	forward

- ✓ (D) Write python code to produce a file in FASTA format named `blast_copies.fasta` containing the sequence for each gene copy listed in (C)

```
from Bio import SeqIO

output = open('blast_copies.fasta', 'w')

blast_result = [
    ('YHR055C', 215064, 214879, '-'), ('YHR055C', 5601, 5786, '+')
]
# sequences = []

for record in SeqIO.parse('exam_data/chr08.fasta', 'fasta'):
    for name, start, end, strand in blast_result:
        subseq = record.seq[start - 1:end]
        print(subseq)
        # SeqIO.write(subseq, output, 'fasta')

    if strand == "-":
        subseq = subseq.reverse_complement()
        print(subseq)
        # SeqIO.write(subseq, output, 'fasta')
```



```
ATGTTTCAGCGAATTAATTAACCTCCAAAATGAAGGTCATGAGTGCCAATGCCAATGTGGTAGCTGCAAAAATGAACAATGCCAAAATCATGTAGCAGCCCAACGCGGTGTAACACCGACG
```

- ✓ (E) Create a multiple sequence alignment from the file you produced in (D).

Note: if you were unable to complete (D), you can still get credit for this question by listing the command you WOULD have run

```
!mafft -auto /content/blast_copies.fasta > blast_copies.MAFFT.fasta
```

- ✓ SIGN HONOR PLEDGE BELOW

On my honor, I have neither received nor given any unauthorized assistance on this examination

KRYSTAL KUANG